

# Task on Modeling

Ignacio Almodóvar & Alejandra Estrada

03/22/2022

## Preprocess

First of all, we load the data.

```
data=read.csv2("drug200.csv",sep = ",",header = TRUE)
```

Then we search for missing values.

```
missingValues=function(data){
  count=0
  a=cbind(lapply(lapply(data, is.na), sum))
  for(i in 1:ncol(data)){
    if(a[i]!=0){
      cat("There are", a[i], "missing values in column ", i,"\n" )
      count=count+1
    }
  }
  if(count==0){
    cat("There are no missing values in this dataset")
  }
}
missingValues(data)
```

```
## There are no missing values in this dataset
```

As we can see, the summary says that all variables except Age are factors. However, analyzing the data the variable “Na\_to\_k” looks like a numeric variable so we must change it.

```
data[,5] %<>% as.numeric()
summary(data)
```


```
##      Age           Sex           BP           Cholesterol
##  Min.   :15.00   Length:200   Length:200   Length:200
##  1st Qu.:31.00   Class :character   Class :character   Class :character
##  Median :45.00   Mode  :character   Mode  :character   Mode  :character
##  Mean   :44.31
##  3rd Qu.:58.00
##  Max.   :74.00
##      Na_to_K      Drug
```

```
## Min.    : 6.269    Length:200
## 1st Qu.:10.445    Class :character
## Median :13.937    Mode  :character
## Mean    :16.084
## 3rd Qu.:19.380
## Max.    :38.247
```

## H2O

We now plot the continuous variables to see if we can find any group evidences for the type of drug

```
ggplot(data,aes(Age,Na_to_K,col=Drug)) + geom_point()
```



Beforehand there is not clear evidence for the differentiation in groups given the age and the Na\_t\_k. However, as can be seen, for the DrugY there is a clear bandwidth for Na\_to\_k being higher than 15.

We are now going to fitt a classification model using h2o package.

```
table(data$Drug)
h2o.init()
data_h2o=as.h2o(data)
resp_data="Drug"
pred_data=setdiff(names(data_h2o), resp_data)

setdiff(names(data_h2o), resp_data)
```

```
data_h2o[, resp_data] <- as.factor(data_h2o[, resp_data])

splits = h2o.splitFrame(data = data_h2o, ratios = 0.8, seed = 42)
train = splits[[1]]
test = splits[[2]]

# Run AutoML
aml_mul = h2o.automl(x = pred_data, y = resp_data, training_frame = train, leaderboard_frame = test,
                    include_algos = c("GLM", "XGBoost", "DeepLearning", "DRF", "GBM", "StackedEnsemble"))
```

If we do leaderboard we obtain the next results:

```
lb_mul <- h2o.get_leaderboard(aml_mul)
head(lb_mul)
```

```
##                                model_id mean_per_class_error
## 1  DeepLearning_grid_1_AutoML_1_20220324_12141_model_1      0.4825397
## 2           GBM_grid_1_AutoML_1_20220324_12141_model_56      0.5111111
## 3 StackedEnsemble_BestOfFamily_4_AutoML_1_20220324_12141      0.5158730
## 4  DeepLearning_grid_1_AutoML_1_20220324_12141_model_12      0.5285714
## 5  DeepLearning_grid_1_AutoML_1_20220324_12141_model_4      0.5333333
## 6           GBM_grid_1_AutoML_1_20220324_12141_model_39      0.5333333
##      logloss      rmse      mse
## 1 1.6372224 0.5839158 0.3409576
## 2 1.0631995 0.5985999 0.3583218
## 3 1.2275490 0.5479127 0.3002083
## 4 0.8657503 0.5753320 0.3310069
## 5 0.8648328 0.5440349 0.2959740
## 6 1.4979195 0.6284823 0.3949900
```

The classification is not really good. The `mean_per_class_error` of the best model of type "DeepLearning", based on fully-connected multilayer artificial neural network, is 0.4015873. As a comparison, the mean-per-class error by “weighted guessing” is:

```
probs <- table(as.matrix(data$Drug))
probs <- probs / sum(probs)
(mean(1 - probs))
```

```
## [1] 0.8
```

And the probability of correct classification by pure chance is:

```
(sum(probs^2))
```

```
## [1] 0.30595
```

We do prediction in the test dataset to see that the prediction is not very good.

```
pred_mul <- h2o.predict(object = aml_mul, newdata = test)
```

```
##      |
```

```
h2o.head(pred_mul)
```

```
##      predict      DrugY      drugA      drugB      drugC      drugX
## 1  drugC 5.806396e-13 5.251812e-02 2.892658e-01 3.781932e-01 2.800229e-01
## 2  drugC 2.193186e-10 5.549356e-02 3.145458e-01 3.420578e-01 2.879028e-01
## 3  DrugY 1.000000e+00 1.569431e-200 1.687176e-164 3.466728e-137 3.504394e-107
## 4  DrugY 9.989732e-01 4.664399e-09 1.191727e-07 1.583093e-09 1.026640e-03
## 5  drugX 4.101541e-17 4.644084e-06 1.650876e-01 1.102436e-02 8.238834e-01
## 6  drugB 5.237041e-15 2.195299e-02 3.760536e-01 3.664320e-01 2.355614e-01
```

And we check the accuracy of the label assignments with the real labels

```
labels_mul <- as.matrix(pred_mul$predict)
table(labels_mul, as.matrix(test$Drug))
```

```
##
## labels_mul drugA drugB drugC drugX DrugY
##      drugA      1      0      0      0      0
##      drugB      0      3      0      2      0
##      drugC      3      0      1      3      0
##      drugX      3      3      1      4      0
##      DrugY      0      0      0      0     14
```

As we expected after observing `h2o.head(pred_mul)` the only label that has been assigned correctly is that of the variable `DrugY`. The rest labels give us very unfavorable results.

So we can conclude that none of these models is very useful since the classification is poor.

To end, we will explain the leader model compare with all AutoML models.

```
ex_mul <- h2o.explain(object = aml_mul, newdata = test)
```

The explainers clearly point to `Na_to_K` being the most relevant predictor. The partial dependence plots are not so useful in this case, we can only conclude that the partial dependence with highest mean response is the partial dependence on “`Na_to_K`” with target = “`DrugY`”.

Finally, we close the `h2o` cluster:

```
h2o.shutdown(prompt = FALSE)
```

## Tidymodels

```
boot_data <- bootstraps(data, times = 10)
analysis(boot_data$splits[[1]] )>% head()
```

```
##      Age Sex      BP Cholesterol Na_to_K Drug
## 103  28  F      LOW           HIGH 13.127 drugC
## 131  70  F NORMAL           HIGH 20.489 DrugY
## 156  49  M      LOW           HIGH 10.537 drugC
##  52  67  M NORMAL      NORMAL 10.898 drugX
##  19  23  M      LOW           HIGH  7.298 drugC
## 110  23  M NORMAL           HIGH 16.850 DrugY
```

## Parsnip

```
library(tidymodels)

# Create an initial split stratifying by the response
set.seed(42)
data_split <- initial_split(data, prop = 0.75)
ames_train <- training(data_split)
ames_test <- testing(data_split)

ames_train$Drug %<>% as.factor()

mnr_spec <- multinom_reg(penalty = 0.1) %>%
  set_engine("nnet")
mnr_spec

## Multinomial Regression Model Specification (classification)
##
## Main Arguments:
##   penalty = 0.1
##
## Computational engine: nnet

mnr_fit <- mnr_spec %>%
  fit(Drug ~ ., data = ames_train)
mnr_fit

## parsnip model object
##
## Fit time: 0ms
## Call:
## nnet::multinom(formula = Drug ~ ., data = data, decay = ~0.1,
##   trace = FALSE)
##
## Coefficients:
##      (Intercept)      Age      SexM      BPLow  BPNORMAL
## drugB -2.2722165  0.152748559 -1.20844566 -1.946016 -1.744401
## drugC  1.6520972  0.004857262 -1.13182828  5.085658 -0.285746
## drugX  0.4929378  0.023586655 -0.89276664  4.297885  5.907929
## DrugY -8.0458019 -0.004643564 -0.07080169  1.089960  1.436271
##      CholesterolNORMAL      Na_to_K
## drugB      0.017088 -0.4091351
## drugC     -2.716356 -0.3381201
## drugX      1.990205 -0.3703790
## DrugY      0.194722  0.6001168
##
## Residual Deviance: 107.2972
## AIC: 163.2972

test_results <- bind_cols(
  dplyr::select(ames_test, "Drug"),
```

```

predict(mnr_fit, ames_test),
predict(mnr_fit, ames_test, type = "prob")
)

table(test_results$Drug, test_results$.pred_class)

```

```

##
##      drugA drugB drugC drugX DrugY
## drugA      4      1      0      0      0
## drugB      0      2      0      0      0
## drugC      0      0      7      0      0
## drugX      0      0      0     12      0
## DrugY      0      0      0      0     24

```

```

mean(test_results$Drug == test_results$.pred_class, na.rm = TRUE)

```

```
## [1] 0.98
```

## Discrim

```

library(discrim)

```

```
## Warning: package 'discrim' was built under R version 4.1.2
```

```

##
## Attaching package: 'discrim'

```

```

## The following object is masked from 'package:dials':
##
##      smoothness

```

*# Fit a Naive Bayes model (which is actually a kernel discriminant analysis done by combining univariate*

```

summary(data)

```

```

##      Age      Sex      BP      Cholesterol
## Min.   :15.00 Length:200 Length:200 Length:200
## 1st Qu.:31.00 Class :character Class :character Class :character
## Median :45.00 Mode  :character Mode  :character Mode  :character
## Mean   :44.31
## 3rd Qu.:58.00
## Max.   :74.00
##      Na_to_K      Drug
## Min.    : 6.269 Length:200
## 1st Qu.:10.445 Class :character
## Median :13.937 Mode  :character
## Mean    :16.084
## 3rd Qu.:19.380
## Max.    :38.247

```

```
data$Drug %<>% as.factor()
nb_mod <- naive_Bayes() %>%
  set_engine("naivebayes") %>%
  fit(Drug ~ ., data = data)
```

```
## Warning: naive_bayes(): Feature BP - zero probabilities are present. Consider
## Laplace smoothing.
```

```
## Warning: naive_bayes(): Feature Cholesterol - zero probabilities are present.
## Consider Laplace smoothing.
```